

The Aricent logo is located in the top right corner. It features the word "Aricent" in a stylized, orange, sans-serif font, with a registered trademark symbol (®) to its upper right. The background of the slide is a dark, abstract composition of glowing, swirling lines in shades of blue, yellow, and orange, creating a sense of motion and energy.

Aricent®

Engineering excellence. **Sourced.**

Linux Device Driver Basics

By

-HEMAKUMAR

Session-2

Character Drivers

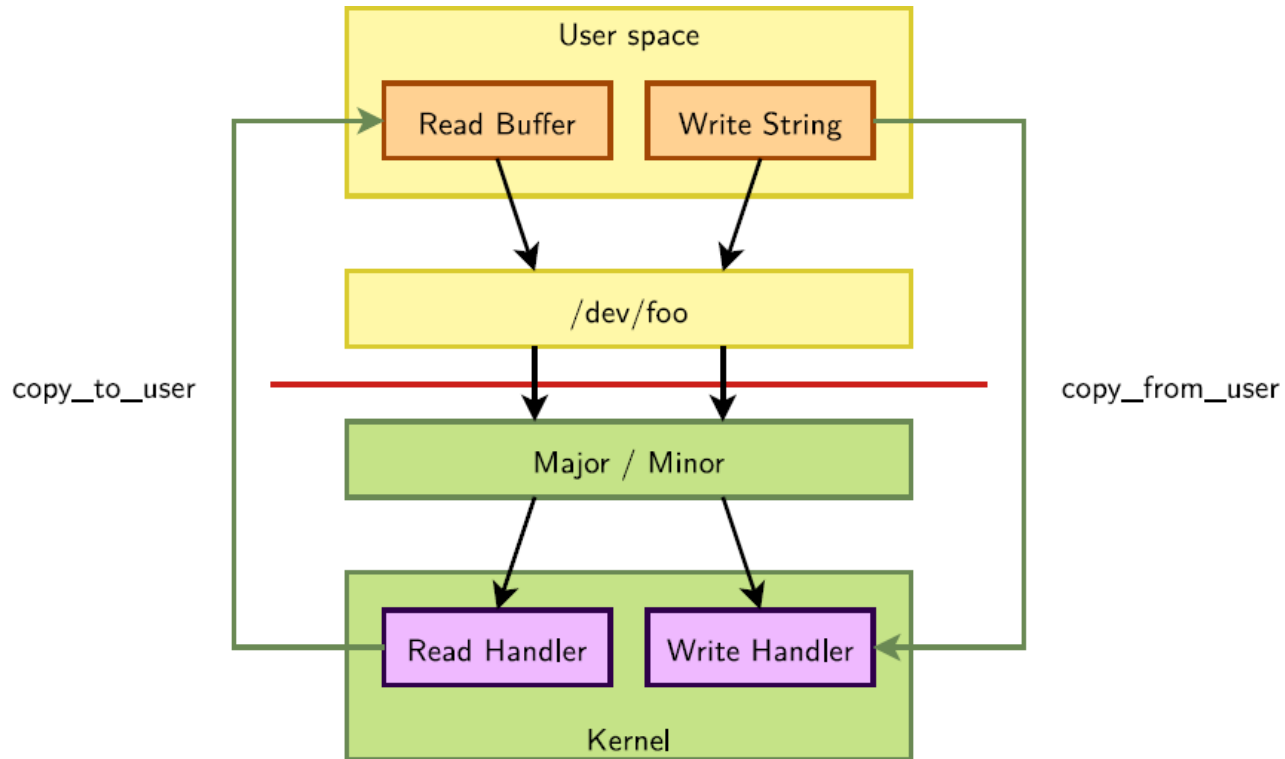
What to Expect?

- ☆ After this session, you would know
 - W's of Character Drivers
 - Major & Minor Numbers
 - Registering & Unregistering Character Driver
 - File Operations of a Character Driver
 - Writing a Character Driver
 - Linux Device Model
 - udev & automatic device creation

Character Driver in Kernel

- From the point of view of an application, a character device is essentially a file.
- The driver of a character device must therefore implement operations that let applications think the device is a file: open, close, read, write, etc.
- In order to achieve this, a character driver must implement the operations described in the struct `file_operations` structure and register them.
- The Linux filesystem layer will ensure that the driver's operations are called when a user space application makes the corresponding system call.

From user space to the kernel: character devices



File operations

- Here are the most important operations for a character driver. All of them are optional.

```
#include <linux/fs.h>
```

```
struct file_operations
```

```
{
```

```
    ssize_t (*read) (struct file *, char_user *, size_t, loff_t *);
```

```
    ssize_t (*write) (struct file *, const char_user *, size_t, loff_t
```

```
    *);    long (*unlocked_ioctl) (struct file *, unsigned int,
```

```
    unsigned long);    int (*mmap) (struct file *, struct
```

```
    vm_area_struct *);
```

```
    int (*open) (struct inode *, struct file
```

```
    *);    int (*release) (struct inode *, struct
```

```
    file *);
```

```
};
```

open() and release()

- `int foo_open(struct inode *i, struct file *f)`

- ▶ Called when user space opens the device file.

- ▶ `struct inode` is a structure that uniquely represents a file in the system (be it a regular file, a directory, a symbolic link, a character or block device)

- ▶ `struct file` is a structure created every time a file is opened. Several file structures can point to the same inode structure.

- ▶ Contains information like the current position, the opening mode, etc.

- ▶ Has a `void *private_data` pointer that one can freely use.

- ▶ A pointer to the file structure is passed to all other operations

- `int foo_release(struct inode *i, struct file *f)`

- ▶ Called when user space closes the file.

read()

- `ssize_t foo_read(struct file *f, char__user *buf, size_t sz, loff_t *off)`
 - ▶ Called when user space uses the `read()` system call on the device.
 - ▶ Must read data from the device, write at most `sz` bytes in the user space buffer `buf`, and update the current position in the file `off`. `f` is a pointer to the same file structure that was passed in the `open()` operation
 - ▶ Must return the number of bytes read. 0 is usually interpreted by userspace as the end of the file.
 - ▶ On UNIX, `read()` operations typically block when there isn't enough data to read from the device

Write()

- `ssize_t foo_write(struct file *f, const char__user *buf, size_t sz, loff_t *off)`

- ▶ Called when user space uses the `write()` system call on the device

- ▶ The opposite of `read`, must read at most `sz` bytes from `buf`, write it to the device, update `off` and return the number of bytes written.

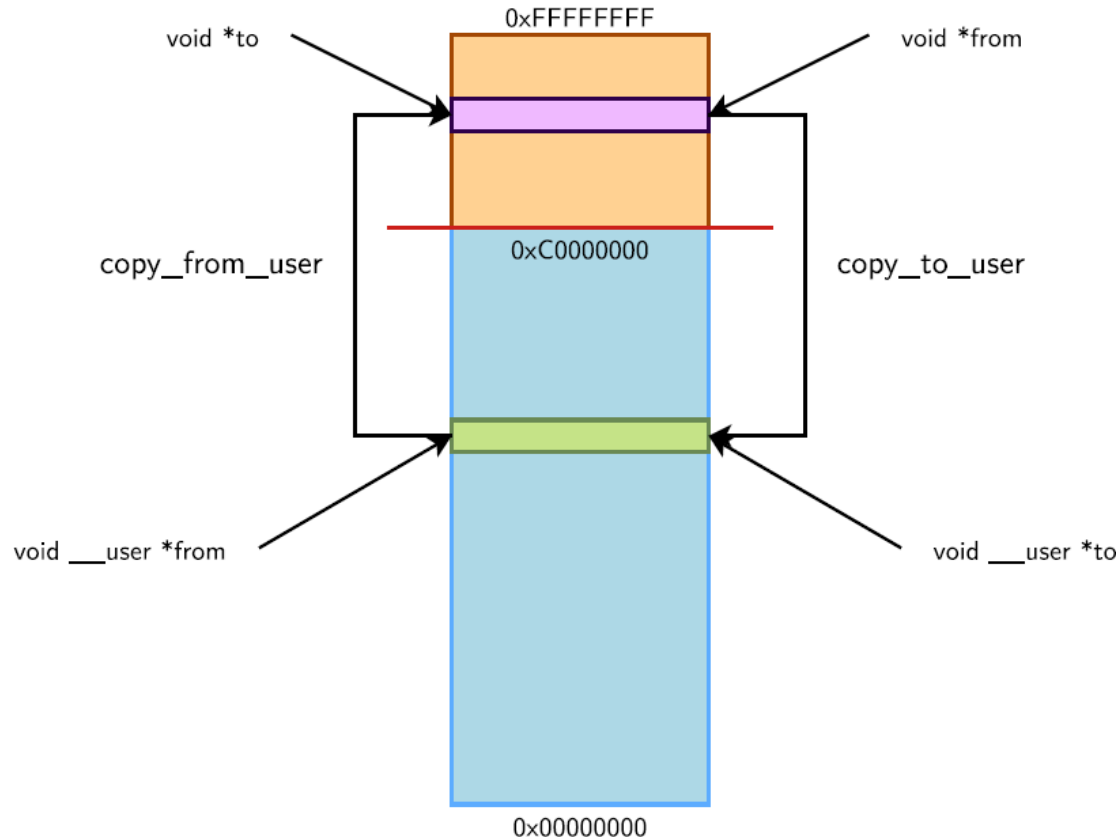
Exchanging data with user space

- Kernel code isn't allowed to directly access user space memory, using `memcpy()` or direct pointer dereferencing.
 - ▶ Doing so does not work on some architectures.
 - ▶ If the address passed by the application was invalid, the application would segfault.
 - ▶ Never trust user space. A malicious application could pass a kernel address which you could overwrite with device data (read case), or which you could dump to the device (write case).
- To keep the kernel code portable, secure, and have proper error handling, your driver must use special kernel functions to exchange data with user space.

Exchanging data with user space.

- A single value
 - ▶ `get_user(v, p);`
 - ▶ The kernel variable `v` gets the value pointed by the user space pointer `p`
 - ▶ `put_user(v, p);`
 - ▶ The value pointed by the user space pointer `p` is set to the contents of the kernel variable `v`.
- A buffer
 - ▶ `unsigned long copy_to_user(void__user *to, const void *from, unsigned long n);`
 - ▶ `unsigned long copy_from_user(void *to, const void__user *from, unsigned long n);`
- The return value must be checked. Zero on success, non-zero on failure. If non-zero, the convention is to return `-EFAULT`.

Exchanging data with user space ..



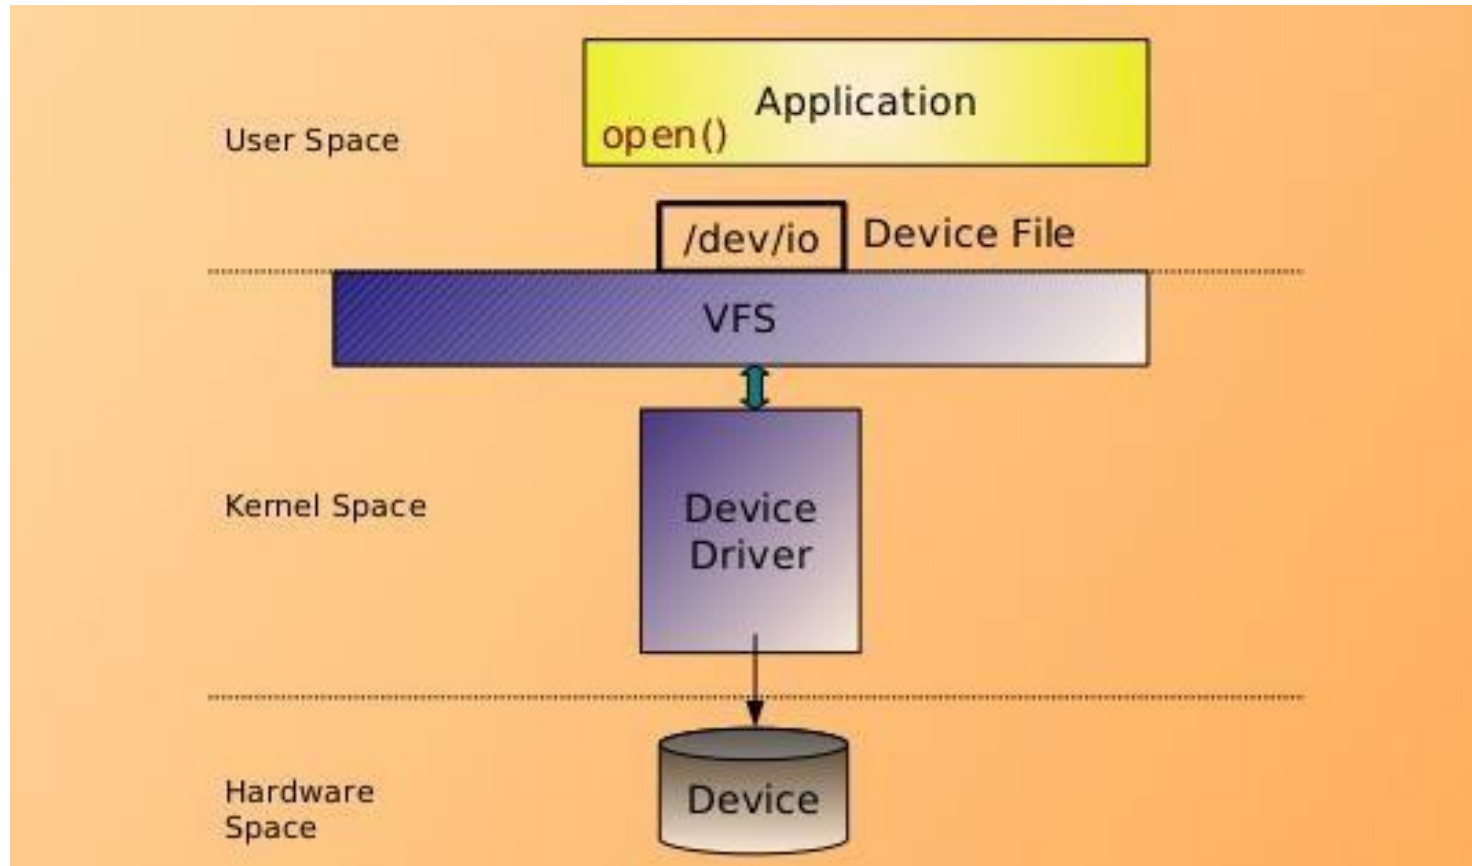
Zero copy access to user memory

- Having to copy data to or from an intermediate kernel buffer can become expensive when the amount of data to transfer is large (video).
- Zero copy options are possible:
 - ▶ `mmap()` system call to allow user space to directly access memory mapped I/O space. See our `mmap()` chapter.
 - ▶ `get_user_pages_fast()` to get a mapping to user pages without having to copy them.

unlocked_ioctl()

- long unlocked_ioctl(struct file *f, unsigned int cmd, unsigned long arg)
 - ▶ Associated to the ioctl() system call.
 - ▶ Called unlocked because it didn't hold the Big Kernel Lock (gone now).
 - ▶ Allows to extend the driver capabilities beyond the limited read/write API.
 - ▶ For example: changing the speed of a serial port, setting video output format, querying a device serial number...
 - ▶ cmd is a number identifying the operation to perform
 - ▶ arg is the optional argument passed as third argument of the ioctl() system call. Can be an integer, an address, etc.
 - ▶ The semantic of cmd and arg is driver-specific.

3 Entities in 3 Spaces



The /dev/null read & write

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t *off)
{
    ...
    return read_cnt;
}

ssize_t my_write(struct file *f, char __user *buf, size_t cnt, loff_t *off)
{
    ...
    return wrote_cnt;
}
```

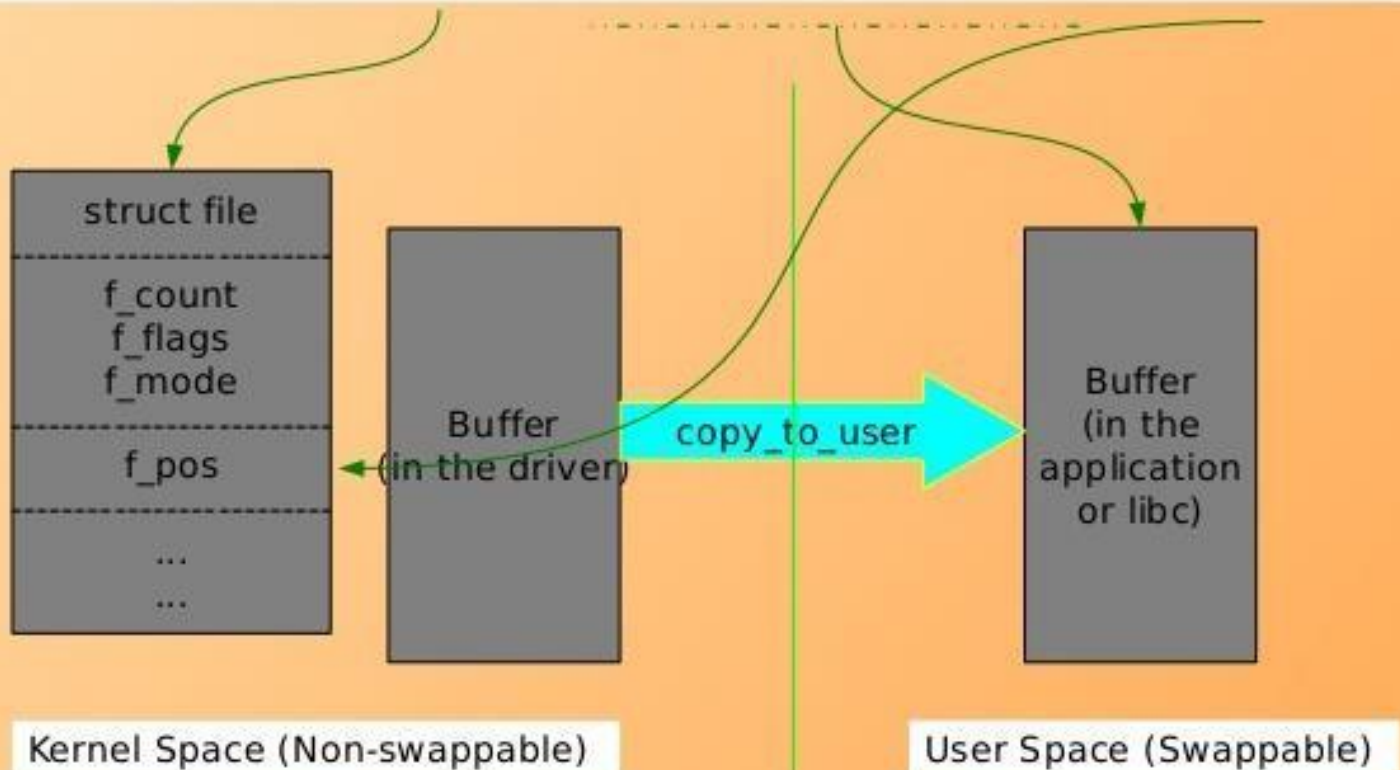

The mem device read

```
#include <asm/uaccess.h>

ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t *off)
{
    ...
    if (copy_to_user(buf, from, cnt) != 0)
    {
        return -EFAULT;
    }
    ...
    return read_cnt;
}
```

The read flow

```
ssize_t my_read(struct file *f, char __user *buf, size_t cnt, loff_t *off)
```



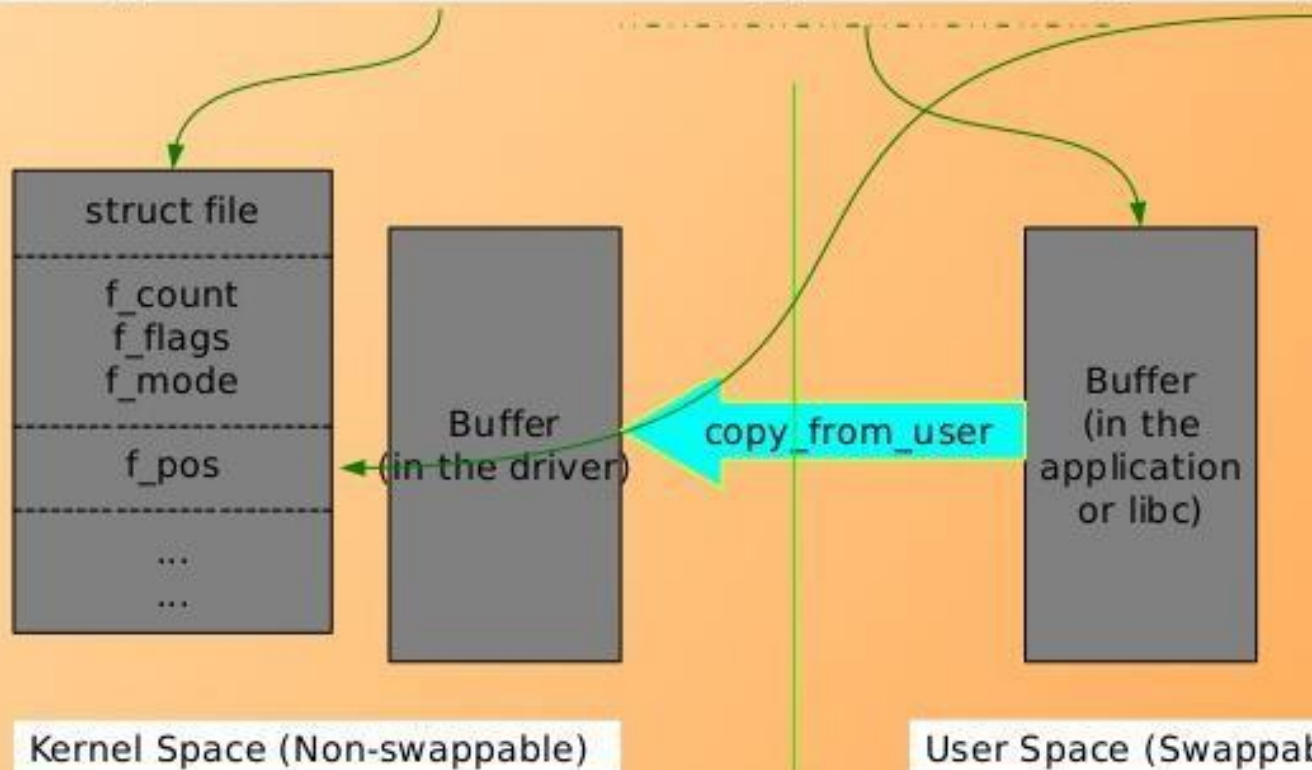
The mem device write

```
#include <asm/uaccess.h>

ssize_t my_write(struct file *f, const char __user *buf, size_t cnt, loff_t *off)
{
    ...
    if (copy_from_user(to, buf, cnt) != 0)
    {
        return -EFAULT;
    }
    ...
    return wrote_cnt;
}
```

The write flow

```
ssize_t my_write(struct file *f, const char __user *buf, size_t cnt, loff_t *off)
```



The mem device write

```
#include <asm/uaccess.h>

ssize_t my_write(struct file *f, const char __user *buf, size_t cnt, loff_t *off)
{
    ...
    if (copy_from_user(to, buf, cnt) != 0)
    {
        return -EFAULT;
    }
    ...
    return wrote_cnt;
}
```

W's of Character Drivers

- ✧ What does “Character” stand for?
- ✧ Look at entries starting with 'c' after
 - `ls -l /dev`
- ✧ Device File Name
 - User Space specific
 - Used by Applications
- ✧ Device File Number
 - Kernel Space specific
 - Used by Kernel Internals as easy for Computation

Major & Minor Number

- ☆ `ls -l /dev`
- ☆ Major is to Category; Minor is to Device
- ☆ Data Structures described in Kernel C in object oriented fashion
- ☆ Type Header: `<linux/types.h>`
 - Type: `dev_t` - 12 bits for major & 20 bits for minor
- ☆ Macro Header: `<linux/kdev_t.h>`
 - `MAJOR(dev_t dev)`
 - `MINOR(dev_t dev)`
 - `MKDEV(int major, int minor)`

Major & Minor Number

Major and Minor Number for a driver:

- Every device driver is identified by a unique ID, which is of 2 parts. Major number and Minor number.
- Major Number: A number indicating which device driver should be used to access a particular device.
- All devices controlled by the same device driver have a common major device number. The minor device numbers are used to distinguish between different devices and their controllers.
- Minor Number: will come in to the picture when there are same kind of multiple devices.

Registering & Unregistering

☆ Registering the Device Driver

- `int register_chrdev_region(dev_t first, unsigned int count, char *name);`
- `int alloc_chrdev_region(dev_t *dev, unsigned int firstminor, unsigned int cnt, char *name);`

☆ Unregistering the Device Driver

- `void unregister_chrdev_region(dev_t first, unsigned int count);`

☆ Header: `<linux/fs.h>`

☆ Kernel Window: `/proc/devices`

Registering the file operations

☆ The Registration

- `int cdev_add(struct cdev *cdev, dev_t num, unsigned int count);`

☆ The Unregistration

- `void cdev_del(struct cdev *cdev);`

☆ Header: `<linux/cdev.h>`

Initialization for Registration

☆ 1st way initialization

- `struct cdev *my_cdev = cdev_alloc();`
- `my_cdev->owner = THIS_MODULE;`
- `my_cdev->ops = &my_fops;`

☆ 2nd way initialization

- `struct cdev my_cdev;`
- `cdev_init(&my_cdev, &my_fops);`

☆ Header: `<linux/cdev.h>`

The file operations

☆ struct file_operations

- struct module owner = THIS_MODULE; /* <linux/module.h> */
- int (*open)(struct inode *, struct file *);
- int (*release)(struct inode *, struct file *);
- ssize_t (*read)(struct file *, char __user *, size_t, loff_t *);
- ssize_t (*write)(struct file *, const char __user *, size_t, loff_t *);
- loff_t (*llseek)(struct file *, loff_t, int);
- int (*unlocked_ioctl)(struct file *, unsigned int, unsigned long);

☆ Header: <linux/fs.h>

The file & inode structures

☆ Important fields of struct file

- mode_t f_mode
- loff_t f_pos
- unsigned int f_flags
- struct file_operations *f_op
- void *private_data

☆ Important fields of struct inode

- unsigned int iminor(struct inode *);
- unsigned int imajor(struct inode *);

The I/O Control API

★ API

- `int (*unlocked_ioctl)(struct file *, unsigned int cmd, unsigned long arg)`

★ Command


➤ Macros

- `_IO, _IOW, _IOR, _IOWR`

➤ Parameters

- type (character) [15:8]
- number (index) [7:0]
- size (param type) [29:16]

dir[31:30]



★ Header: `<linux/ioctl.h>` →...→ `<asm-generic/ioctl.h>`

SIMPLE CHARACTER DRIVER - Implementation of LCD driver

CHARACTER DRIVER of LCD driver:

```
#include <linux/modules.h>
#include <linux/version.h>
#include <linux/kernel.h>
#include <linux/fs.h>
#include <linux/init.h>
#include <linux/uaccess.h>

#define LCD_MAJOR 254 static
unsigned char inuse = 0; static
int nbytes;

int mylcd_open(struct inode * inode, struct file * filp)
{
    printk("my_LCD _ open Invoked \n");
    if(inuse)
        return EBUSY;
    inuse = 1;
    return 0;
}

int mylcd_release(struct inode * inode, struct file * filp)
{
    inuse = 0;
    return 0;
}
```

SIMPLE CHARACTER DRIVER - Implementation of LCD driver..

```
Size_t mylcd_write(struct file * file, const char *buf, size_t count, loff_t *f_pos)
```

```
{    char data[10];  
    nbytes = copy_from_user(data, buf, count);  
    printk("\n data = %s", data);  
    return nbytes;  
}
```

```
Static struct file_operations fops = {
```

```
    write:  mylcd_write,  
    open :  mylcd_open,  
    release: mylcd_release,
```

```
};
```

```
Int mylcd_init(void)
```

```
{  
    int result = 0;  
    inuse = 0;  
    result = register_chrdev(LCD_MAJOR, "mylcd", &fops);  
    return 0;  
}
```

```
Void mylcd_exit(void)
```

```
{  
    unregister_chrdev(LCD_MAJOR, "mylcd");  
}
```

```
MODULE_DESCRIPTION("TEST LCD DRIVER");
```

```
MODULE_AUTHOR("MY_TEAM");
```


Test App to SIMPLE CHARACTER DRIVER - Implementation of LCD driver...

Test application in U-space:

```
#include <stdio.h>
#include <sys/types.h>
#include <fcntl.h>

Main()
{
    int fd, n;
    char buf[12] = "Hello World";
    fd = open("/dev/mylcd", O_RDWR);
    Printf("\n fd = %d", fd);
    getchar();
    n = write(fd, buf, 10);
    printf(" No. of bytes read %d\n", n);
}
```

SIMPLE CHARACTER DRIVER - Implementation of LCD driver....

Compilation & Execution of above driver:

```
[root@testkerdrv]# make
[root@testkerdrv]# insmod mylcd.ko
[root@testkerdrv]# cat /proc/devices | more -- to check driver inserted successfully or
not.
char devices:
1 mem
4 /dev/vc/0
254 mylcd
256 pcidev
[root@testkerdrv]# mknod /dev/mylcd c 254 0
[root@test]# ./a.out
fd = 3;    -- open call got executed
[root@testkerdrv]# dmesg
mylcd open invoked.
data = hello world.
```

Ioctl() example: Kernel side

```
static long phantom_ioctl(struct file *file, unsigned int cmd, unsigned long arg)
{
    struct phm_reg r;
    void __user *argp = (void __user *)arg;

    switch (cmd) {
    case PHN_SET_REG:
        if (copy_from_user(&r, argp,
            sizeof(r))) return -EFAULT;
        /* Do something
        */ break;
    case PHN_GET_REG:
        if (copy_to_user(argp, &r,
            sizeof(r))) return -EFAULT;
        /* Do something
        */ break;
    default:
        return -ENOTTY;
    }
    return 0; }
```

Ioctl() Example: Application Side

```
int main(void)
{
    int fd, ret;
    struct phm_reg reg;
    fd =
    open("/dev/phantom");
    assert(fd > 0);
    reg.field1 = 42;
    reg.field2 = 67;
    ret = ioctl(fd, PHN_SET_REG, &
    reg); assert(ret == 0);
    return 0;
}
```

Linux Device Model (LDM)

- ✧ struct kobject - <linux/kobject.h>
 - ✧ kref object
 - ✧ Pointer to kset, the parent object
 - ✧ kobj_type, type describing the kobject
- ✧ kobject instantiation → sysfs representation
- ✧ Parent object guides the entries under /sys/
 - ✧ bus – the physical buses
 - ✧ class – the device categories
 - ✧ device – the actual devices

udev & LDM

- ★ Daemon: udevd
- ★ Configuration: /etc/udev/udev.conf
- ★ Rules: /etc/udev/rules.d/
- ★ Utility: udevinfo [-a] [-p <device_path>]
- ★ Receives uevent on a change in /sys
- ★ Accordingly, updates /dev &/or
- ★ Performs the appropriate action for
 - Hotplug
 - Microcode / Firmware Download
 - Module Autoload

Device Model & Classes

- ✧ Latest way to create dynamic devices
 - Create or Get the appropriate device category
 - Create the desired device under that category
- ✧ Class Operations
 - `struct class *class_create(struct module *owner, char *name);`
 - `void class_destroy(struct class *cl);`
- ✧ Device into & out of Class
 - `struct class_device *device_create(struct class *cl, NULL, dev_t devnum, NULL, const char *fmt, ...);`
 - `void device_destroy(struct class *cl, dev_t devnum);`

What all have we learnt?

- ☆ W's of Character Drivers
- ☆ Major & Minor Numbers
- ☆ Registering & Unregistering Character Driver
- ☆ File Operations of a Character Driver
- ☆ Writing a Character Driver
- ☆ Linux Device Model
- ☆ udev & automatic device creation


Any Queries?

Books for Ref..

Books:

- Understanding the Linux Kernel, D. P. Bovet and M. Cesati, O'Reilly & Associates, 2000.
- Linux Core Kernel – Commentary, In-Depth Code Annotation, S. Maxwell, Coriolis Open Press, 1999.
- TheLinux Kernel, Version 0.8-3, D. ARusling, 1998.
- Linux Kernel Internals, 2nd edition, M. Beck et al., Addison-Wesley, 1998.
- Linux Kernel, R. Card et al., John Wiley & Sons, 1998.
- Linux Device Drivers, 3rd Edition, Jonathan Corbet, Alessandro Rubini, and Greg Kroah-Hartman Published by O'Reilly Media, Inc., 1005

Questions ??



Thank you!