

Peripheral Component Interconnect (PCI) Drivers

What to Expect?

- ☆ PCI Subsystem

- Coding Details of a PCI Device Driver
- Browsing Sample Code

PCI Prologue

- ✧ Developed by Intel as a replacement for ISA
 - ✧ More of a processor (specific) bus for x86 architectures
- ✧ Replacing Features
 - ✧ Plug n Play
 - ✧ Platform independent
 - ✧ Fast
- ✧ Hardware Addressing
 - ✧ Through South Bridge (PCI Controller)
 - ✧ Buses (upto 256) → Devices (upto 32) → Functions (upto 8)

Kernel Aspect

- ☆ Organized as
 - PCI Core (driver of the PCI Controller)
 - PCI Drivers (device drivers)
- ☆ Kernel Windows for PCI
 - `/proc/bus/pci/devices` (Try cat)
 - `/sys/bus/pci/devices` (Try tree)
- ☆ Command: `lspci`

Dynamic Device Configuration

★ What Configuration?

- Device Base Address(es)
 - Register Space(s)
 - Memory(s)
- Interrupt Vectors
- ...

★ Done during Bootup by

- BIOS, Or
- PCI Core (Bootloader or Kernel)

★ Configured by the time, driver comes up

★ But, where is this configuration stored?

PCI Device Configuration Space

	0x0										0xF
0x00	Vendor ID	Device ID	Command Register	Status Register	Rev ID	Class Code	Cache Line	Lat Timer	Header Type	BIST	
0x10	Base Address 0		Base Address 1		Base Address 2		Base Address 3				
0x20	Base Address 4		Base Address 5		CardBus CIS pointer		Subsystem Vendor ID		Subsystem Device ID		
0x30	Expansion ROM Base Address		Reserved				IRQ Line	IRQ Pin	Min Gnt	Max Lat	

PCI Device Config Space Access

★ Header: <linux/pci.h>

★ APIs

- `int pci_read_config_byte(struct pci_dev *dev, int where, u8 *val);`
- `int pci_read_config_word(struct pci_dev *dev, int where, u16 *val);`
- `int pci_read_config_dword(struct pci_dev *dev, int where, u32 *val);`
- `int pci_write_config_byte(struct pci_dev *dev, int where, u8 *val);`
- `int pci_write_config_word(struct pci_dev *dev, int where, u16 *val);`
- `int pci_write_config_dword(struct pci_dev *dev, int where, u32 *val);`

NB `struct pci_dev *dev` would be available as parameter to `probe()`

Standard PCI Device Access API

- ★ Upto 6 Memory or I/O regions

- ★ Header: <linux/pci.h>

- ★ APIs

 - unsigned long pci_resource_start(dev, bar);

 - unsigned long pci_resource_len(dev, bar);

 - unsigned long pci_resource_flags(dev, bar);

- ★ Flags

 - IORESOURCE_IO

 - IORESOURCE_MEM

 - IORESOURCE_PREFETCH

PCI Driver Registration

- ✧ `int pci_register_driver(struct pci_driver *drv);`
- ✧ `int pci_unregister_driver(struct pci_driver *drv);`
- ✧ `struct pci_driver`
 - `const char *name`
 - `const struct pci_dev_id *id_table;`
 - `PCI_DEVICE(vendor, device);`
 - `PCI_DEVICE_CLASS(dev_class, dev_class_mask);`
 - `int (*probe)(pci_dev, id_table);`
 - `void (*remove)(pci_dev);`
- ✧ Header: `<linux/pci.h>`

Typical 'probe' Function

```
int probe(struct pci_dev *dev, struct pci_dev_id *id)
{
    /* Enable the PCI Device */
    pci_enable_device(dev);
    ...
    /* Acquire PCI Device's regions */
    pci_request_regions(dev, "label");
    ...
    /* Possibly map I/Os */
    ...
    pci_set_drvdata(dev, <pvt_data_ptr>); // Optional
    ...
    /* Register the vertical */
    ...
    return 0; /* Claimed. Negative for not Claimed */
}
```


Typical 'remove' Function

```
void remove(struct pci_dev *dev)
{
    /* Unregister the vertical */
    ...
    pci_set_drvdata(dev, NULL); // Optional
    ...
    /* Possibly unmap I/Os */
    ...
    /* Release PCI Device's regions */
    pci_release_regions(dev);
    ...
    /* Disable the PCI Device */
    pci_disable_device(dev);
}
```

Old-style PCI Probing & Getting

- ☆ `struct pci_dev *pci_find_(vendor, device, from);`
- ☆ `struct pci_dev *pci_get_device(v, d, from);`
- ☆ `struct pci_dev *pci_get_subsys(v, d, ssv, ssd, f);`
- ☆ `struct pci_dev *pci_get_slot(bus, devfn);`
- ☆ `pci_dev_put(pci_dev);`

DMA Mapping




★ APIs

- ▶ `dma_addr_t pci_map_single(struct pci_dev *, void *, size_t, int dir);`
- ▶ `void pci_unmap_single(struct pci_dev *, dma_addr_t, size_t, int dir);`
- ▶ Directions
 - `PCI_DMA_BIDIRECTIONAL`
 - `PCI_DMA_TODEVICE`
 - `PCI_DMA_FROMDEVICE`

★ Header: `<linux/pci.h>`

DMA Allocation



★ APIs

- `void *pci_alloc_consistent(struct pci_dev *, size_t, dma_addr_t *);`
- `void pci_free_consistent(struct pci_dev *, size_t, void *, dma_addr_t);`
- `int pci_set_dma_mask(struct pci_dev *, u64 mask);`

★ Header: `<linux/pci.h>`

What all have we learnt?

- ☆ PCI Subsystem

- Coding Details of a PCI Device Driver
- Browsing Sample Code

- ☆ DMA with PCI

Any Queries?