

# Introduction to PCIe driver

03.12.2020

# Introduction to PCIe driver

**03.12.2020**



# Agenda

- 1** PCI Architecture
- 2** PCI Topology
- 3** PCIe Basics and Concepts
- 4** PCIe Limitation
- 5** PCIe Performance
- 6** PCI Address Space

# Agenda

**7** PCI Memory and I/O Spaces

**8** PCIe Layers

**9** PCI Linux Commands

**10** PCIe Flow Control

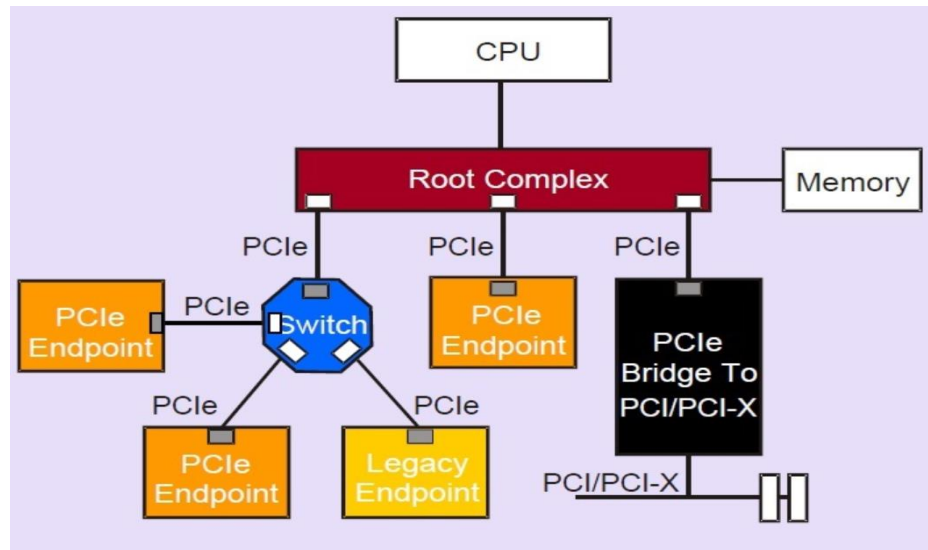
**11** PCI Driver Example

**12** PCI Applications

**13** PCIe 6.0

## PCI Architecture

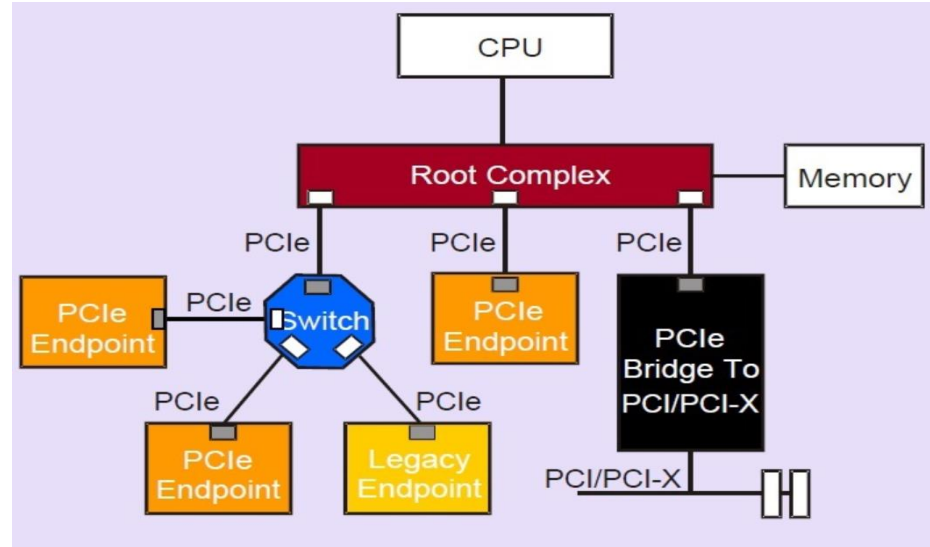
- PCI Express is based on point-to-point topology with separate serial links connecting every device to the root complex (host).
- PCI Express bus link supports full-duplex communication between any two endpoints.
- The PCI Express link between two devices can vary in size from one to 32 lanes.
- The PCI Express standard defines link widths of x1, x2, x4, x8, x12, x16 and x32.
- PCI Express preserves backward compatibility with PCI; legacy PCI system software can detect and configure newer PCI Express devices without explicit support for the PCI Express standard.



# PCI Architecture & Topology

## PCIe Components:

- Root Complex
- End Points
- PCIe to PCI(X) Bridge
- Requester
- Completer
- Port
- Switch




1. Root Complex is interface between the CPU and the PCIe buses may contain several components.
2. End Pointes can be Wire, Wireless cards, FPGAs.
3. Requester is who initiate the commands (read, write).
4. Complete is who accepts the commands and process that.
5. Ports where endpoints/switches connected
6. Switches are used to connect the multiple endpoints


# PCIe Basics and Concepts

Dual Simplex Point to Point serial connection

- Serial, point-to-point, Low Voltage diff signals
- Independent transmit and receive sides
- Scalable Link Widths (x1,x2,x4,x8,x12,x16,x32)
- Scalable Link Speeds (2.5 to 32 GT/s [64 GT/s for PCIe 6.0])
- Packet based transaction protocol
- Virtual Channels, Traffic classes
- End to end CRC
- Flow Control



x4 Link Example



## PCI Limitation

- Limited host pin-count.
- PCI(X) uses a shared parallel bus architecture, in which the PCI host and all devices share a common set of address, data and control lines, So it is limited to one master at a time.
- Parallel bus required more wires between host and endpoint where PCIe uses Serial bus required less wires.
- Data Rate : PCI
  - PCI 133 MB/s (32-bit at 33 MHz)
  - PCI 266 MB/s (32-bit at 66 MHz)
  - PCI-X 266 MB/s (64-bit at 33 MHz)
  - PCI-X 533 MB/s (64-bit at 66 MHz)



# PCIe Performance

PCIe Bandwidth:

Data Rate (GB/s) aprx	Link Width				
	x1	x2	x4	x8	x16
PCIe 1.x	0.25	0.5	1	2	4
PCIe 2.x	0.5	1	2	4	8
PCIe 3.0	1	2	4	8	16
PCIe 4.0	2	4	8	16	32
PCIe 5.0	4	8	16	32	64
PCIe 6.0	8	16	32	64	128

# PCI Configuration Address Space

- Contains 256 or 4K bytes of basic device information.
- The first 64 bytes (00h –3Fh) make up the standard configuration header, including PCI ID, i.e. vendor ID and device ID registers, to identify the device.
- The remaining 192 bytes (40h –FFh) represent user-definable configuration space.

## Vendor ID:

This 16 bit value identifies the manufacturer of the function.

## Device ID:

This 16 bit value is assigned by manufacturer to identify the type of the function.

31		16 15		0
Device ID		Vendor ID		00h
Status		Command		04h
Class Code			Revision ID	08h
BIST	Header Type	Lat. Timer	Cache Line S.	0Ch
Base Address Registers				10h
				14h
				18h
				1Ch
				20h
Cardbus CIS Pointer				24h
				28h
				2Ch
Subsystem ID		Subsystem Vendor ID		30h
Expansion ROM Base Address				34h
Reserved			Cap. Pointer	38h
Reserved				3Ch
Max Lat.	Min Gnt.	Interrupt Pin	Interrupt Line	

# PCI Configuration Address Space

## Revision ID:

This 8 bit value is assigned by manufacturer to identify the revision of the function.

## Class Code:

Defines the basic function of the functions

- 01 Mass storage controller
- 02 Network controller
- 03 Display controller (etc)

## Base Address Registers:

These registers are used to determine and allocate the type, amount and location of PCI I/O and PCI memory space that the device can use.

The **Subsystem ID** and the **Subsystem Vendor ID** differentiate specific model (such as an add-in card).

While the Vendor ID is that of

the chipset manufacturer, the Subsystem Vendor ID is that of the card manufacturer.

31

16

15

0

Device ID		Vendor ID		00h
Status		Command		04h
Class Code			Revision ID	08h
BIST	Header Type	Lat. Timer	Cache Line S.	0Ch
Base Address Registers				10h
				14h
				18h
				1Ch
				20h
				24h
Cardbus CIS Pointer				28h
Subsystem ID		Subsystem Vendor ID		2Ch
Expansion ROM Base Address				30h
Reserved			Cap. Pointer	34h
Reserved				38h
Max Lat.	Min Gnt.	Interrupt Pin	Interrupt Line	3Ch

# PCI Configuration Address Space

## Header Type:

Type 0 for endpoints

Type 1 for RC, switches and bridges.

- Type 0 configuration requests are sent only to the device for which they are intended.
- Type 1 configuration requests are sent to switches/bridges on the way; the last one before the actual target device will convert it to type 0.
- A device that receives a type 0 request knows that this request is intended for this particular device, so it uses the bus/device number fields in the request to find out what its own bus address is.

31		16 15		0		
Device ID		Vendor ID				00h
Status		Command				04h
Class Code				Revision ID		08h
BIST	Header Type	Lat. Timer		Cache Line S.		0Ch
Base Address Registers						10h
						14h
						18h
						1Ch
						20h
						24h
Cardbus CIS Pointer						28h
Subsystem ID		Subsystem Vendor ID				2Ch
Expansion ROM Base Address						30h
Reserved				Cap. Pointer		34h
Reserved						38h
Max Lat.	Min Gnt.	Interrupt Pin		Interrupt Line		3Ch

# PCI Config Space Read/Write Functions in Linux

```
int pci_read_config_byte(const struct pci_dev *dev, int addr, u8 *val);
```

```
int pci_read_config_word(const struct pci_dev *dev, int addr, u16 *val);
```

```
int pci_read_config_dword(const struct pci_dev *dev, int addr, u32 *val);
```

```
int pci_write_config_byte(const struct pci_dev *dev, int addr, u8 val);
```

```
int pci_write_config_word(const struct pci_dev *dev, int addr, u16 val);
```

```
int pci_write_config_dword(const struct pci_dev *dev, int addr, u32 val);
```

# PCI Config Space Read/write Commands in Linux

Accessing configuration space from user space PCI configuration space access from user space is possible via sysfs.

```
# lspci | grep VGA
```

**01:00.0 VGA compatible controller: ATI Technologies Inc RV370 [Sapphire X550 Silent]**

```
# hexdump /sys/bus/pci/devices/0000\:00\:00.0/config
```

```
00000000 8086 7192 0006 0200 0003 0600 0000 0000
```

```
0000010 0000 0000 0000 0000 0000 0000 0000 0000
```

# PCI Memory and I/O Spaces

Memory space is used by most everything else – it's the general-purpose address space

- The PCI spec recommends that a device use memory space, even if it is a peripheral
- An agent can request between 16 bytes and 2GB of memory space. The PCI spec recommends that an agent use at least 4kB of memory space, to reduce the width of the agent's address decoder
- IO space is where basic PC peripherals (keyboard, serial port, etc.) are mapped
- The PCI spec allows an agent to request 4 bytes to 2GB of I/O space

# PCIe Layers

Application Data transferred via packets:

Transaction Layer Packet (TLP)

Data Link Layer Packet (DLLP)

TLP:

Memory Read Unlocked

Memory Read Locked

Memory Write

IO Read/Write

Conf Read

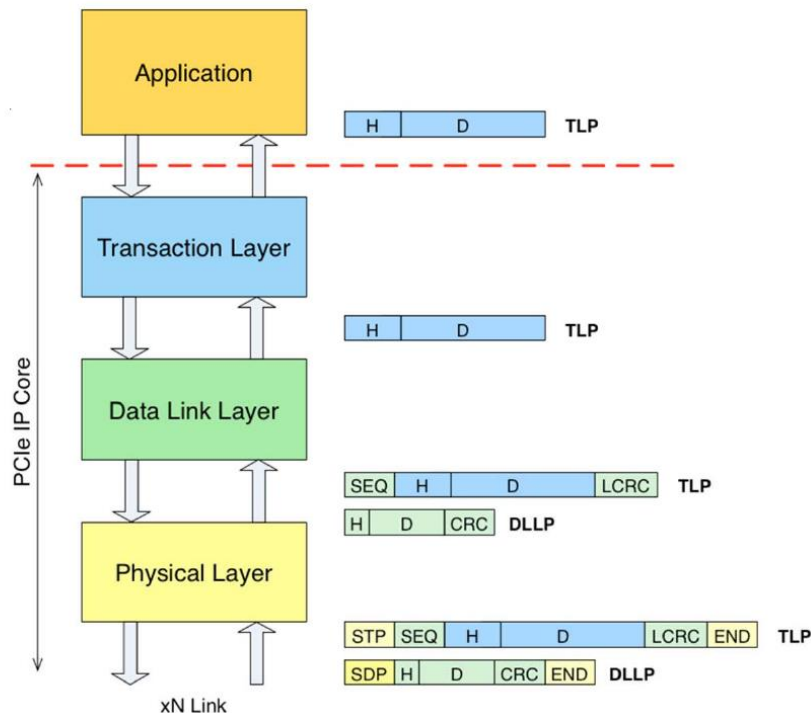
Conf Write

Completion (locked, Unlocked)

Completion with Data (Locked, Unlocked)

Msg Req

Msg Req with Data



DLLP:

TLP ACK/NAK

Power Management

Link Flow Control



# PCIe Bus Enumeration

- Bus enumeration is performed by attempting to read the vendor ID and device ID register for each combination of bus number and device number at the device's function #0.
- After a new bridge is detected, a new bus number is defined, and device enumeration restarts at device number zero
- If a valid vendor ID is returned from bus 0, device 0, function 0, this indicates that the device is implemented and contains at least one function. Looking for remaining functions (1-7). Proceed to next device.
- If a value of FFFFh were returned as the Vendor ID, this would indicate that function 0 is not implemented in device 0 and the enumeration software (OS/device driver) will proceed for next device.
- OS/device driver software attempts to read the Vendor ID from function 0 in each of the 32 possible devices on bus.
- If the Vendor ID succeeds, Software writes all ones to its BARs and reads back the device's requested memory size and operating system will program the memory-mapped and I/O port addresses into the device's BAR configuration register.

# PCIe Interrupts

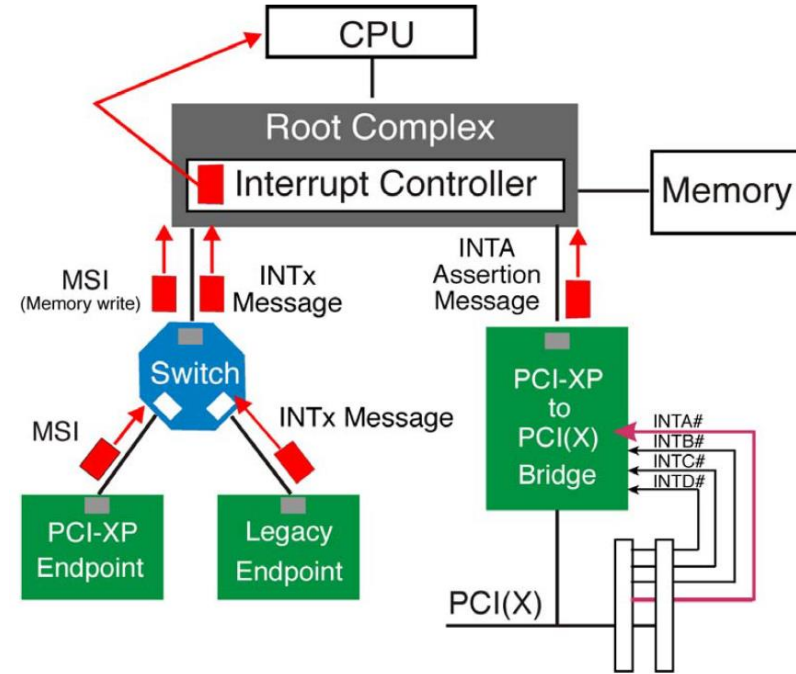
**Types of interrupts:** Legacy Interrupt, MSI (Message Signaled Interrupts) or MSI-X depending on their design requirements.

## Legacy PCI Interrupt Delivery:

- Legacy functions use one of the interrupt lines (INTA#, INTB#, INTC#, INTD#) to signal an interrupt.
- INTx# signal is asserted to request interrupt service.
- PCIe defines in-band messages that act as virtual INTx# wires which target the interrupt controller located typical within the Root Complex.

## Limitation of Legacy PCI Interrupts:

- 4 Interrupt pins needs to be shared across all endpoints/bridges/switches. Need to parse the source of the interrupt.



# PCIe Interrupts

## MSI/MSI-X:


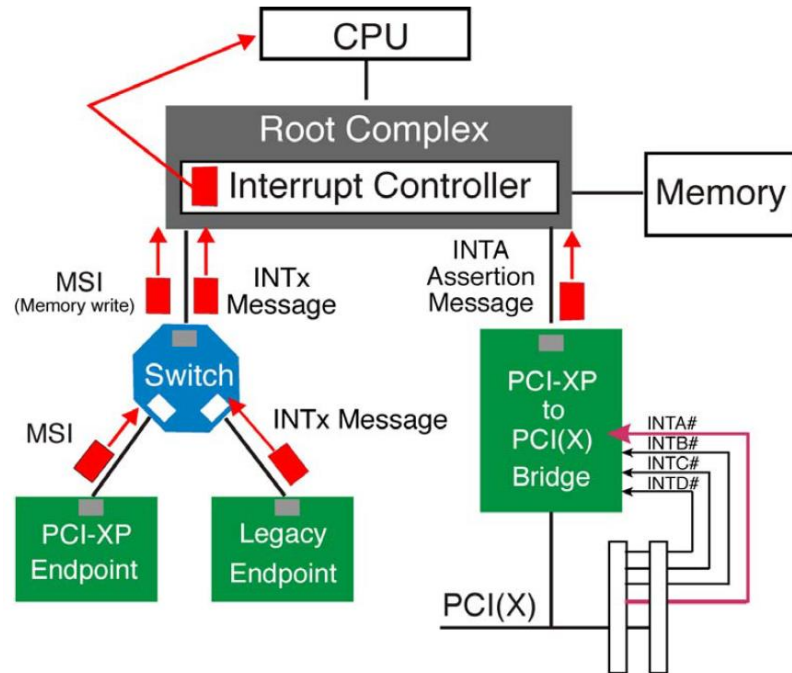
- No need any sideband signals.
- Simple memory Write transactions.
- MSI supports 32 interrupt vectors
- MSI-X supports 2048 interrupt vectors
- MSI can only be distinguished from other memory writes by the address locations they target, which are reserved by the system for interrupt delivery.

## MSI Capability Register:

31	16	15	8	7	0	
Message Control Register				Pointer to Next ID	Capability ID = 05h	Dword 0
Least-Significant 32-bits of Message Address Register				0	0	Dword 1
Most-Significant 32-bits of Message Address Register				Message Data Register		Dword 2
						Dword 3

31	16	15	8	7	0	
Message Control Register				Pointer to Next ID	Capability ID = 05h	Dword 0
Message Address Register				0	0	Dword 1
				Message Data Register		Dword 2



## PCI Linux commands

# lspci - display all the device information

- 00:00.0 Host bridge: Intel Corporation 5500 I/O Hub to ESI Port (rev 13)
- 00:01.0 PCI bridge: Intel Corporation 5520/5500/X58 I/O Hub PCI Express Root Port 1 (rev 13)
- 00:09.0 PCI bridge: Intel Corporation 7500/5520/5500/X58 I/O Hub PCI Express Root Port 9 (rev 13)
- 00:14.0 PIC: Intel Corporation 7500/5520/5500/X58 I/O Hub System Management Registers (rev 13)

# lspci -t

```
-[0000:00]--+-00.0
  +-01.0-[01]--+-00.0
    |      \-00.1
  +-03.0-[02]--+-00.0
    |      \-00.1
  +-07.0-[04]
```

## PCI Linux commands

# lspci -v

03:00.0 RAID bus controller: LSI Logic / Symbios Logic MegaRAID SAS 2108 [Liberator] (rev 05)

Subsystem: Dell PERC H700 Integrated

Flags: bus master, fast devsel, latency 0, IRQ 16

I/O ports at fc00 [size=256]

Memory at df1bc000 (64-bit, non-prefetchable) [size=16K]

Memory at df1c0000 (64-bit, non-prefetchable) [size=256K]

Expansion ROM at df100000 [disabled] [size=256K]

Capabilities: [50] Power Management version 3

Capabilities: [68] Express Endpoint, MSI 00

Capabilities: [d0] Vital Product Data

Capabilities: [a8] MSI: Enable- Count=1/1 Maskable- 64bit+

Capabilities: [c0] MSI-X: Enable+ Count=15 Masked-

Capabilities: [100] Advanced Error Reporting

Kernel driver in use: megaraid\_sas

Kernel modules: megaraid\_sas

# PCI Linux commands

## # setpci

A utility for querying and configuring PCI devices

“**setpci -s 12:3.4 3c.l=1,2,3**” writes longword 1 to register 3c, 2 to register 3d and 3 to register 3e of device at bus 12, slot 3, function 4

## # setpci --dumppregs

cap pos w name

00 W VENDOR\_ID

02 W DEVICE\_ID


04 W COMMAND

06 W STATUS

08 B REVISION (etc)

# PCIe Flow Control

Flow control guarantees that transmitters will never send Transaction Layer Packets (TLPs) that the receiver can't accept. This prevents receive buffer over-runs and eliminates the need for inefficient disconnects, retries, and wait-states on the link. Receiver sends Flow Control Packets (FCP) which are a type of DLLP (Data Link Layer Packet) to provide the transmitter with credits so that it can transmit packets to the receiver.



# Linux PCI Driver Example

```
static struct pci_device_id ids[] =  
{  
    { PCI_DEVICE(VENDOR_ID, DEVICE_ID) }  
};
```

```
static struct pci_driver pci_drv =  
{  
    .name = DEVICE_NAME,  
    .id_table = ids,  
    .probe = device_probe,  
    .remove = device_remove,  
};
```



# Linux PCI Driver Example

```
static int __init device_init(void)
{
    int status = 0;

    //Register PCI device. This will match the Device Vendor ID from table and call the probe function
    status = pci_register_driver(&pci_drv);
    if (status) {
        printk("PCI driver failed to register, rval = %d\n", status);
    }
    return status;
}

static void __exit device_exit(void)
{
    pci_unregister_driver(&pci_drv);
    return;
}
```

# Linux PCI Driver Example

```
static int device_probe(struct pci_dev *dev, const struct pci_device_id *id)
{
    unsigned int base0;
    unsigned long memflag = 0;
    unsigned long memlen = 0;
    unsigned char *bar0_addr;

    //Read PCI BAR0 address (Shared mem address)
    pci_read_config_dword(dev, PCI_BASE_ADDRESS_0, &base0);

    //Returns bus start address for a given PCI region
    memstart = pci_resource_start(dev, BAR0);
    //Returns the byte length of a PCI region
    memlen = pci_resource_len(dev, BAR0);

    //check memlen before requesting the memory region
    if (request_mem_region(memstart, memlen, dev->dev.kobj.name) == NULL)
    {
        return -ENOMEM;
    }
}
```

# Linux PCI Driver Example

```
//Enable PCI device
pci_enable_device(dev);

//Get the virtual address of the BAR0
bar0_addr = (unsigned char *) ioremap(base0, memlen);

return 0;
}
```

```
static void device_remove(struct pci_dev *dev)
{
    //Disable PCI device
    pci_disable_device(dev);
}
```

# PCI Applications

PCI Express operates in consumer, server, and industrial applications, as a motherboard-level interconnect (to link motherboard-mounted peripherals), a passive backplane interconnect and as an expansion card interface for add-in boards.

## External GPUs

- 3D Graphic Cards
- Soundcards

## Network

- 10G or multigigabit cards
- 802.11ax WiFi cards

## Enterprise Storage

- RAID Cards
- SSD drives
- Flash memories

## PCIe 6.0

Planned to be released in 2021

- Bandwidth is expected to increase to 64 GT/s, yielding 126 GB/s in each direction and 256GB/s(bidirectional) in a 16-lane configuration
- 4-level pulse-amplitude modulation(PAM-4) signaling will be used
- Reed—Solomon forward error correction(RS-FEC), a predictive error correction.
- Maintains backward compatibility to previous generations
- PCI-Express 6.0 is also likely to be useful in self-driving systems, the industrial IoT, and any system that juggles multiple sensors to combine input from multiple peripherals into a cohesive whole

# REVISION HISTORY

Rev No.	Date	Description of Change	Author	Review and Approved By
1.0	03-12-2020	Initial Draft	Rajeshkumar R	Suresh K

alTRAN

Part of Capgemini 