

Process Management

On Linux

Prayas Mohanty (Red Hat Certified Instructor)

Red Hat Certification ID: 100-005-594

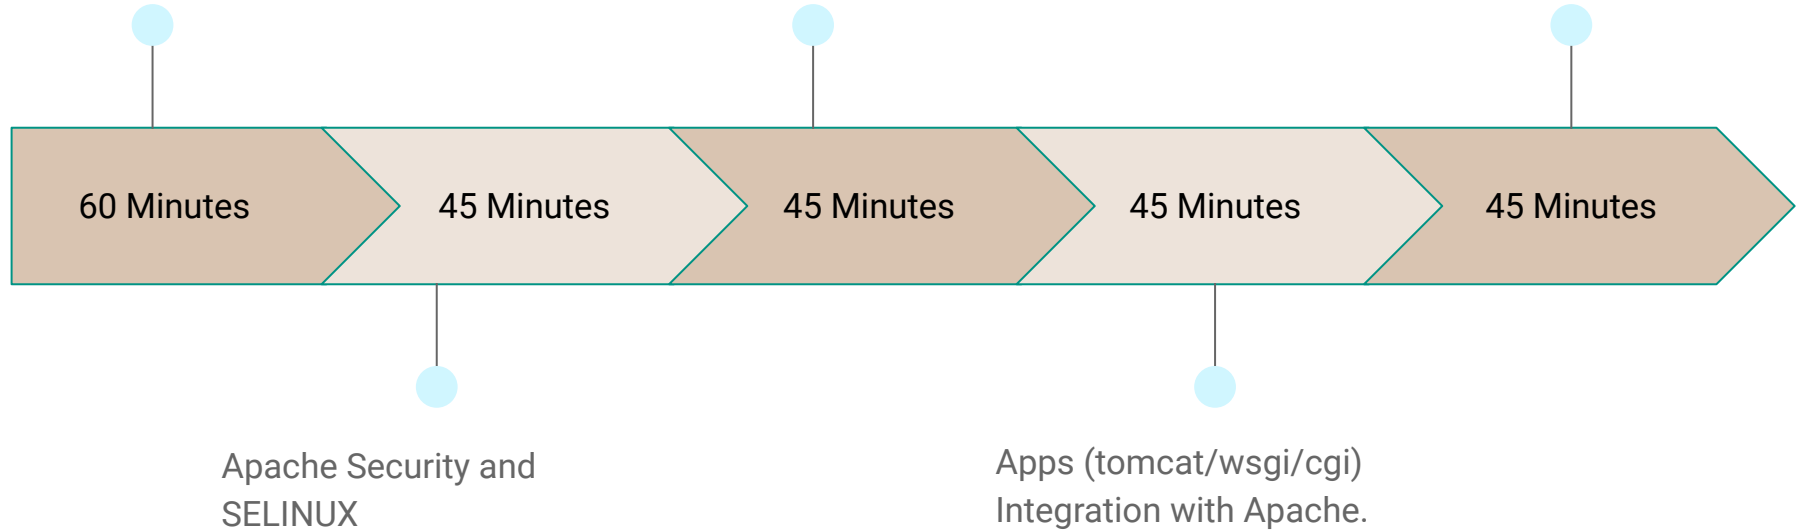
Objective

- What is the concepts of process management.
 - Process creation
 - Parent – Child Relation in Process
 - Process Termination & the Role of wait.
 - What are zombie and orphan process.
- What Linux system calls – fork(), wait().... do
- How to write programs using system call to:
 - create processes
 - extract exit codes using macros
 - handle zombie processes etc.
- What are process specific coding guidelines.

Securing Apache using SSL
module

Optimizing Apache using
MPM Modules

Automate Web Server
Deployment using Ansible.



Prerequisite of Participants

- Having Knowledge on Linux Platform
- Having familiarity with Linux Command line
- Basic Knowledge on vi Editor
- Proper Knowledge on C programming
- Basic Knowledge on Processes

What is Process Management

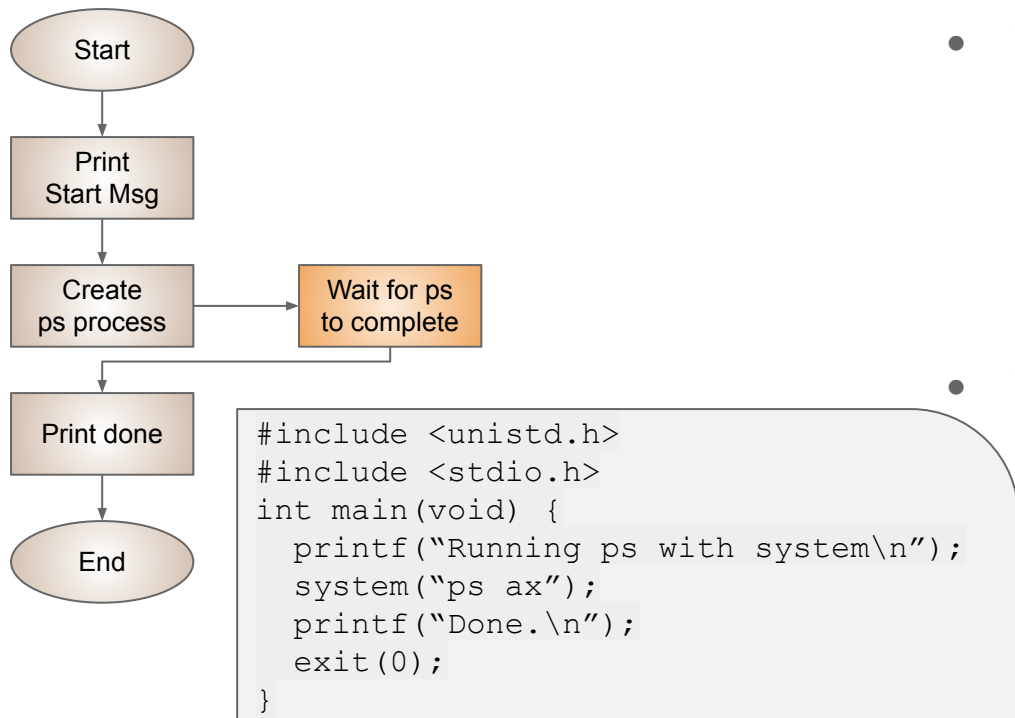
- Process management involves various tasks like creation, scheduling, termination of processes.
-

Process Creation

- Everytime you execute a command, Linux create a process for your command.
 - Commands are executables stored in a directory specified by PATH variable.
 - Executing a command loads the instructions from the executable file to the text part of the address space after creating the necessary data structure for a Process with a PID.
- You can cause a program to run from inside another program and thereby create a new process by using the system library function.

```
#include <unistd.h>
#include <stdio.h>
int main(void) {
    printf("Running ps with system\n");
    system("ps ax");
    printf("Done.\n");
    exit(0);
}
```

How that program works

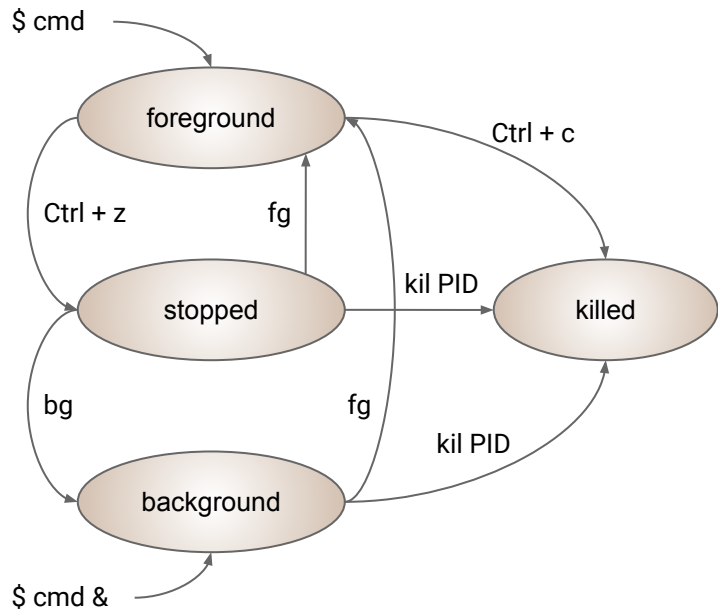


- The program calls system function with the string "ps ax" , which executes the ps program.
 - In this case the running instance of the program calling the system function is call the **Parent** process.
 - The running instance of the ps program is called the **Child** process.
- The program returns from the call to system when the ps command has finished.
 - As a default, Parents wait for Child process to gracefully complete to start execution of it's next instruction(wait to ready).

How to Control the Process



- When you run a command / create a process, Your process may need some input and produce some output / errors.
- By default, all the inputs are fetched from the user through keyboard and the output, and error are written to the Monitor.
- This combination of Monitor and keyboard are technically referred as terminal (tty) which is sometime physical and most often virtual.
- All these processes which are being attached to a terminal can also be called jobs for that terminal and mostly running in foreground.
- The foreground jobs are attached to the terminal session's lifecycle.
 - If the terminal is closed, the respective job gets terminated.
- Most of the server programs do not tolerate this terminal dependency, so they prefer to run on background.

Create a process to run in background



- The job in the background continues to run irrespective of the terminal from which it is being initiated.
- You either run a job in the background or move an existing running foreground job in the to the background.
- All the inputs and outputs are supposed to be read and written to a file.
 - if you put a job in the background which prints back to the terminal you might still get to see the output of the job in the terminal.
- This kind of background job can be terminated using kill command & process id.

Create a process to run in background



```
#include <unistd.h>
#include <stdio.h>
int main(void) {
    printf("Running ps with system\n");
    system("ps ax &");
    printf("Done.\n");
    exit(0);
}
```

- The program call to system returns as soon as the ps program is started as it was to start ps command in background.
- The program then prints Done, and exits before the ps command has had a chance to finish all of its output.
- This kind of process behavior can be quite confusing for users.
- To make good use of processes, you need finer control over their actions.

Food for thought

Write program to start 10 background count program for 10 minutes and show the result?

Replacing the Process Image

- In general, using `system` is a far from ideal way to start other processes.
 - A much better way of invoking programs is to use `system` call instated.
- You can use `exec` functions to “hand off” execution of your program to another.
- An `exec` function replaces the current process with a new process specified by the path & file argument.
- There is a whole family of related functions grouped under the `exec` heading.
 - `int execl(const char *path, const char *arg0, ..., (char *)0);`
 - `int execlp(const char *file, const char *arg0, ..., (char *)0);`
 - `int execlx(const char *path, const char *arg0, ..., (char *)0, char *const envp[]);`
 - `int execv(const char *path, char *const argv[]);`
 - `int execvp(const char *file, char *const argv[]);`
 - `int execve(const char *path, char *const argv[], char *const envp[]);`

Exec Function Family

- Execl Family
 - Take all argument Independently
- Execv Family
 - Take all argument as a single array of string
- Both Family
 - P take the path parameter
 - E need to define own (Customize) environment variable.

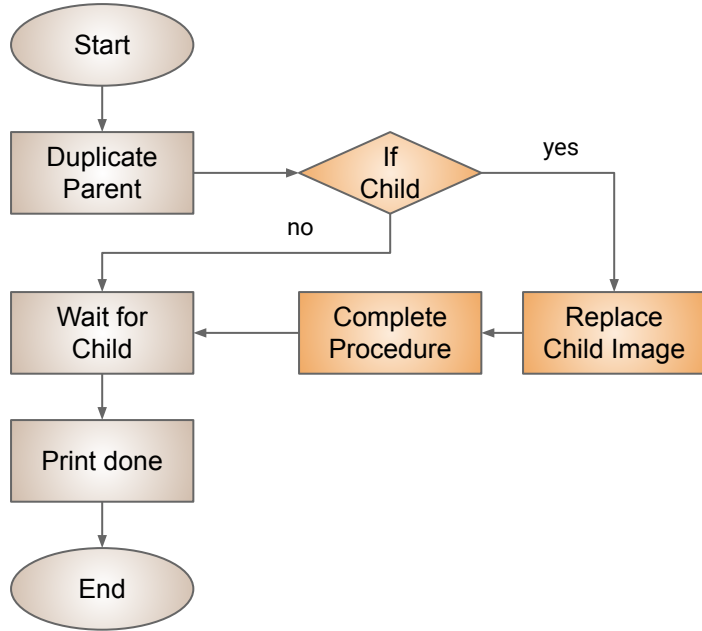
```
#include <unistd.h>
#include <stdio.h>
int main(void) {
    printf("Running ps with exelp\n");
    execlp("ps", "ps", "ax", 0);
    printf("Done.\n");
    exit(0);
}
```

```
#include <unistd.h>
#include <stdio.h>
int main(void) {
    char *const ps_argv[] = {"ps", "ax", 0};
    char *const ps_envp[] = {"PATH=/bin:/usr/bin", "TERM=console",
0};
    printf("Running ps with execlve\n");
    execve("/bin/ps", ps_argv, ps_envp);
    printf("Done.\n");
    exit(0);
}
```

Two major difference between system & exec

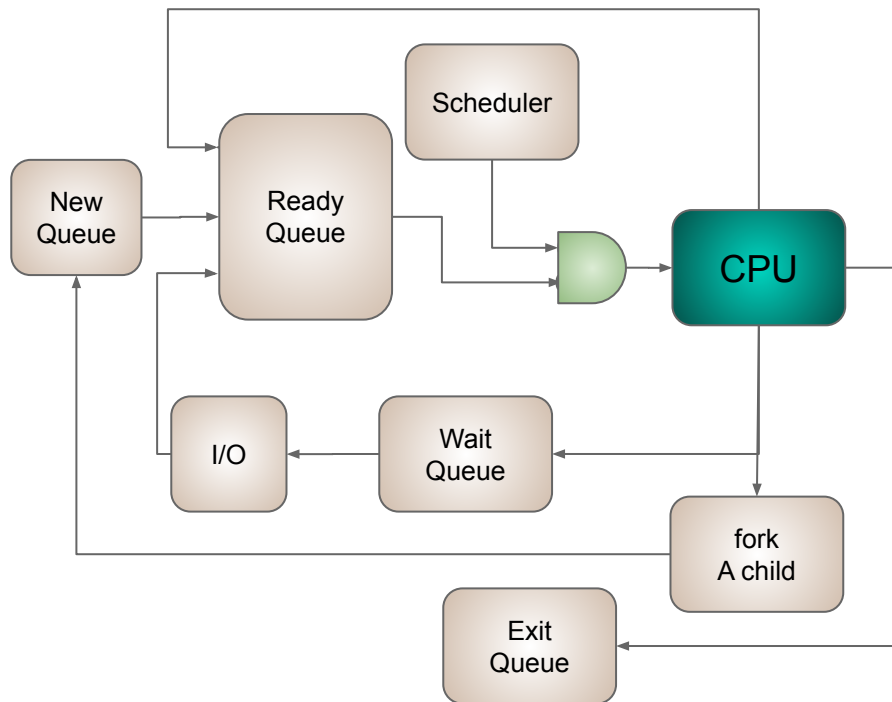
- **There is no reference to a current process in the ps output.**
 - The program prints its first message and then calls exec function family , which replace the current process image with ps and start execute new instruction from start, which is now ps.
 - The PID of the new process is the same as the original, along with the parent PID and nice value.
 - In effect, all that has happened is that the running program has started to execute new code from a new executable file specified in the call to exec .
- **When you run this program, you get the usual ps output, but no Done.**
 - Because of the image replacement, when ps finishes, it return back to it's parent (shell) that provide you the shell prompt.
 - As the control doesn't return to the original program , so the second message doesn't get printed.
- This kind of process behavior can be quite confusing for users.

Best Practice for multiprocessing application



- Best Practice for multiprocessing application is to create a new child process for all new task.
- The easiest way to create a child process is to use a system call fork.
 - fork duplicate the parent image into a new process treated as the child process.
- In child process use the exec family system call to replace the child image to parent image.
- As parent wait for child to complete and then start executing the remaining procedure.

Creating Process with System Call (fork)



- This program runs as two processes creating a new process as Duplicate of parent.
 - The process calling fork() is called the parent.
 - The new process is called the child of the parent.
- Fork is usually a system call, implemented in the kernel.
- Fork is the primary (and historically, only) method of process creation on Unix-like operating systems.

```
#include<stdio.h>
#include<unistd.h>
int main()
{
    pid_t newpid;
    newpid = fork();
    printf("Hello, I have PID %d\n", (int) getpid());
    return 0;
}
```


What about concurrency

```
include<unistd.h>
#include<stdio.h>
int main() {
    pid_t pid; char *message; int n;
    puts("fork program starting\n");
    pid = fork();
    switch(pid) {
        case 0:
            message = "This is the child";
            n = 5;
            break;
        default:
            message = "This is the parent";
            n = 3;
            break;
    }
    for(; n > 0; n--) {
        puts(message);
        sleep(1);
    }
    return 0;
}
```

- The child and parent processes are executed concurrently.

The wait system call

```
#include<sys/wait.h>
#include<unistd.h>
#include<stdio.h>
int main() {
    pid_t pid,childpid; char *message; int n,status;
    puts("fork program starting\n");
    pid = fork();
    switch(pid) {
        case 0:
            message = "This is the child";
            n = 5;
            break;
        default:
            message = "This is the parent";
            n = 3;
            childpid = wait(&status);
            break;
    }
    for(; n > 0; n--) {
        puts(message);
        sleep(1);
    }
    return 0;
}
```

- A call to wait() blocks the calling process until one of its child processes exits or a signal is received.
 - You must include <sys/wait.h>
 - It return the PID of the terminated child.
 - The status of the child can be captured as an integer by passing a pointer to the wait call.
- If you want to wait for a specific child use waitpid() instated.
 - Like wait it also return the PID of terminated child.
 - Along with status it take 2 more argument as option and the PID

The exit system call

- When a process terminates it executes an `exit()` system call, either directly, or indirectly via library code.
- This leaves an exit status value (typically an integer) in the PCB of the terminated process for the parent process to read later.
- You must include `<stdlib.h>`
-

Zombie and orphan process

```
#include<unistd.h>
#include<stdio.h>
int main() {
    pid_t pid; char *message; int n =2;
    puts("fork program starting\n");
    pid = fork();
    switch(pid)
    {
        case 0:
            message = "This is the child";
            break;
        default:
            message = "This is the parent";
            getchar();
            break;
    }
    for(; n > 0; n--) {
        puts(message);
        sleep(1);
    }
    return 0;
}
```

- When a child process terminates, an association with its parent survives until the parent in turn either terminates normally or calls wait .
- Although no longer active, the child process is still in the system because its exit code needs to be stored in case the parent subsequently calls wait.
- This state of the process is known as defunct, or a **Zombie process**.
- If the parent terminates abnormally, the child process automatically gets the process with PID 1 as parent and known as **Orphan process**.
- A zombie orphan process is the last thing one expect to have in his system.

Error Handling

```
#include<unistd.h>
#include<stdio.h>
int main() {
    pid_t pid; char *message; int n =2;
    puts("fork program starting\n");
    pid = fork();
    switch(pid)
    {
        case 0:
            message = "This is the child";
            break;
        default:
            message = "This is the parent";
            getchar();
            break;
    }
    for(; n > 0; n--) {
        puts(message);
        sleep(1);
    }
    return 0;
}
```

- The status information collected from wait system call allows the parent process to determine the exit status of the child process.
- You can interpret the status information using macros defined in sys/wait.h , shown in the following table.

Error Handling

```
#include<unistd.h>
#include<stdio.h>
int main() {
    pid_t pid,child;char *message;
    int n=4,exitcode=0,status;
    puts("fork program starting\n");
    pid = fork();
    switch(pid) {
        case -1:
            perror("fork failed");
            exit 1;
        case 0:
            message = "This is the child";
            exitcode = 37;
            break;
        default:
            message = "This is the parent";
            getchar();
            break;
    }
}
```

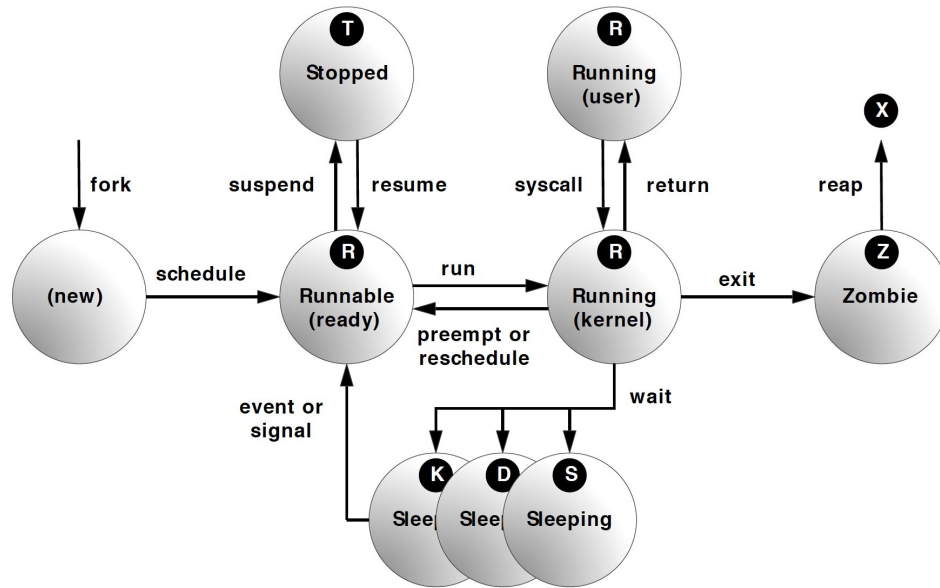
- The status information collected from wait system call allows the parent process to determine the exit status of the child process.
- You can interpret the status information using macros defined in sys/wait.h , shown in the

```
for(; n > 0; n--) {
    puts(message);
    sleep(1);
}
if (pid != 0) {
    child = wait(&status);
    printf("Child has finished:PID=%d\n",child);
    if(WIFEXITED(status))
        printf("Child exit code %d\n", WEXITSTATUS(status));
    else
        printf("Child terminated abnormally\n");
}
exit(exitcode);
```

Listing / controlling Processes through ps & top

- View Process information with ps
 - Shows process from the current terminal by default
 - a includes process on all terminal
 - x includes process not attached to terminal
 - u process owner information
- The Linux version of ps supports three option formats :
 - UNIX(POSIX) which start with a -
 - BSD options, where - is optional
 - GNU long options, which support --
- For example, ps -aux is not the same as ps aux.
- Use top command to interactively list and control the current process.

Process status as per ps output



- $R \rightarrow \text{TASK_RUNNING}$
- $S \rightarrow \text{TASK_INTERRUPTIBLE}$
- $D \rightarrow \text{TASK_UNINTERRUPTIBLE}$
- $K \rightarrow \text{TASK_KILLABLE}$
- $T \rightarrow \text{TASK_STOPPED}$
- $Z \rightarrow \text{EXIT_ZOMBIE}$
- $X \rightarrow \text{EXIT_DEAD}$

What are process specific coding guidelines

- Handle errors after system call.
- Wait for child exit.
- Extract exit code using `WEXITSTATUS()`.
- Avoid zombies.

Useful Links

-

Summary

- What is the concepts of process management.
 - Process creation
 - Parent – Child Relation in Process
 - Process Termination & the Role of wait.
 - What are zombie and orphan process.
- What Linux system calls – fork(), wait().... do
- How to write programs using system call to:
 - create processes
 - extract exit codes using macros
 - handle zombie processes etc.
- What are process specific coding guidelines.

Any Questions ?

Thank you!